

Semantic information inside sentence embeddings

Anonymous ACL submission

Abstract

This is abstract

1 Introduction

2 Related work

Costra corpus (Barančíková and Bojar, 2020)

SBERT (Reimers and Gurevych, 2019)

Doc2vec (Le and Mikolov, 2014)

3 Experiments

3.1 Costra corpus

3.2 Embeddings

We focus mainly on two sets of embeddings: doc2vec predicted by Doc2vec model trained on 1M of Czech sentences and sbert predicted by pretrained SBERT multilingual model. Additionally we also experimented with supervised method using the same SBERT model — we call this embedding sbert_supervised.

To obtain sbert_supervised we splitted the corpus into 5 folds. For each fold, the model trained on the remaining 4 folds and produced embeddings for the one remaining fold. This gave us 5 sets of disjunct embeddings, each trained on slightly different data. For training we used softmax loss, which classifies sentence and its corresponding seed sentence into the set of all possible transformations. This gives the model incentive to produce embeddings that are easily distinguishable for different transformations.

3.3 Tasks

3.3.1 Probing tasks

We designed a set of probing tasks which aim to test whether the embeddings contain semantic information while ignoring how exactly it is encoded. More specifically we ask a classifier to predict the transformation used to obtain the given embedding.

The motivation for these tasks is that if the embedding contains semantic information (in our case the type of transformation used), it would be beneficial for the classifier to find and use this information. This would boost its scores and make it stand out.

We test two settings: *no-context* and *contextualized*. In no-context setting the classifier receives only the sentence embedding itself and test which (if any) transformations can be reliably determined without the knowledge of the seed sentence. In the contextualized setting we pass both the sentence embedding and the embedding of the corresponding seed sentence. This enables the classifier to examine the relationships between the transformed and seed sentences. The results are presented in Section 4.1, the complete method is described in detail in Appendix B.

4 Results

4.1 Probing for transformation prediction

Besides the already described embeddings in Section 3.2, we also add their versions with randomly permuted labels within sentences with the same seed sentence (named mixup_by_seed_) and several types bag-of-words (*BOW*) vectors. The mixup_by_seed_ embeddings serve as a comparative random baseline. We experimented with two types of BOW vectors — pure bow vector having ones at positions corresponding to words included in the given sentence (named bow) and TF-IDF weighted bow vectors (named tfidf). Additionally we tried to limit the amount of words registered in order to decrease the dimensionality of the vectors, however this failed to increase the classifiers' scores and therefore we omit them in the rest of this section. Dimensions of all embeddings are presented in Table 1.

We employ several types of classifiers: Random Forests (RandomForestClassifier), Sup-

Embedding	Dimension	Sparseness
doc2vec	512	0%
sbert	384	0%
bow	8374	99.89%
tfidf	8374	99.89%

Table 1: Dimensionality and sparseness as the percentage of zeros of used embeddings.

port Vector Machine (SVC) and K-Nearest neighbours (KNeighboursClassifier).

In summary, our entire pipeline is as follows:

1. For each embedding, we consider only a random sample with uniform distribution of transformations and ignore the embeddings of the seed sentences.
2. For each embedding and classifier we run a 5-fold cross-validated grid search over several parameters, which gives us the best model for given embedding.
3. We train each embedding on the classifier which performed the best using cross-validation.

More details about the pipeline can be found in Appendix B.

4.1.1 Contextualized setting

We tried two methods of contextualization: *concatenating* the two embeddings and *subtracting* the seed embedding from the sentence embedding. Since the method using difference performed better, we omit the concatenation results.

We report the grid search results in Figure 1. The best performing embeddings were those produced by SBERT. Supervised SBERT embeddings differ significantly from the unsupervised ones only in superior performance of KNeighboursClassifier. This suggests better local structure in terms of the sentence transformation.

Moreover the BOW embeddings surpassed doc2vec embedding, proving that for prediction of transformation, lexical features are more valuable than any information extracted by our trained Doc2vec model.

In Figure 2 we show the performance of best classifiers per each label for the most performant dense and lexical embeddings: sbert and tfidf. We can identify four transformations as being well-predicted having around 0.80 f1-score. However




Figure 1: Grid search F1 macro averaged scores of contextualized embeddings. The reported scores of sbert_supervised are the mean scores of all SBERT supervised embeddings.

the powerful processing of SBERT offered significant advantage to pure lexical information only for two of those: ‘past’ and ‘future’. We see even bigger difference for ‘opposite meaning’, even if it reached just 0.4 f1-score.




Figure 2: Comparison of F1 score per each transformation of best classifiers for sbert and tfidf embedding in contextualized setting.

In order to examine the most confused transformations we summed confusion matrices of the top 3 most performant embeddings. The resulting matrix is shown in Figure 3. The most confused pairs of transformations are ‘minimal change’ with

‘different meaning’ and ‘nonsense’ with ‘different meaning’, both in around 40%. Interestingly ‘paraphrase’ gets recognised almost never, with ‘simple sentence’, ‘formal sentence’, ‘nonsense’ and ‘minimal change’ having accuracy below 20%.

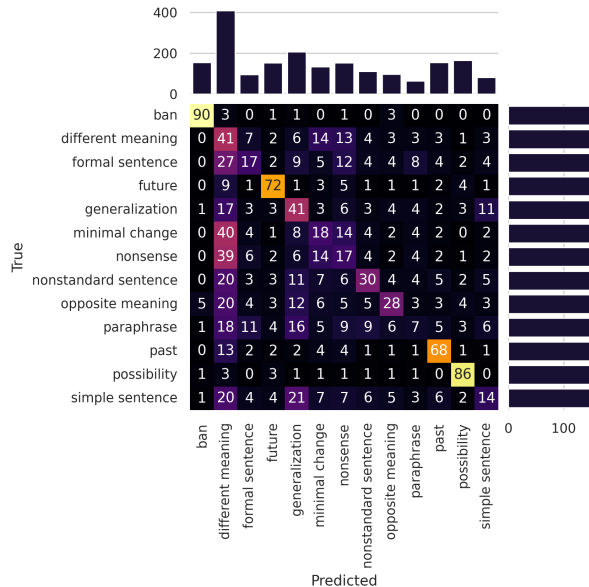


Figure 3: Summed confusion matrices of best classifiers for sbert, sbert_supervised and tfidf embeddings. The confusion matrix is row-normalized, whereas the displayed true and predicted distributions are not.

4.1.2 No-context setting

Grid search results for no-context setting are displayed in Figure 4. Comparing these results to those in contextualized setting, we see large drop of performance for sbert embedding — almost by 50%. This shows how much the classifiers relied on the information contained in the embedding of seed sentence. This drop is not as pronounced for sbert_supervised embedding, since the model producing it is optimized for good scores in precisely this type of task.¹ On the other hand, the BOW embeddings performed almost as well, suggesting they were not really able to take advantage of the contextualized setting.

In Figure 5 we compare performance of best classifiers for sbert and tfidf embedding. Whereas the most recognised transformations are the same as in the contextualized setting, the classifiers performance significantly drops in all cases except for the ‘ban’ and the ‘possibility’ transformations embedded with TF-IDF. This suggests that ‘ban’ and

¹For this reason, we avoid comparing sbert_supervised to other embeddings from this point onwards.




Figure 4: Grid search F1 macro averaged scores of sentence embeddings with no context. The reported scores of sbert_supervised are the mean scores of all SBERT supervised embeddings.

‘possibility’ transformations can be easily identified from pure *lexical* features. While SBERT fails to encode this information into its embedding, it performs even better when the embeddings of a sentence and its corresponding seed sentence are put into one context. This phenomena is pronounced mainly for ‘different meaning’ and ‘past’ transformations, where the inclusion of seed sentence’s embedding on the input is responsible for over 80% and 60% of F1 score of the classifier, as can be seen in Figure 6. Overall SBERT’s embeddings benefit from contextualization more than TF-IDF vectors.

4.2 Evaluating Sentence Embeddings with Gaussian Mixture

This section presents an unsupervised approach for evaluating the separability of sentence embeddings. We measure label separability pairwise using a Gaussian Mixture model and calculate the F-score of the unsupervised clustered labels with the true labels.

The Gaussian Mixture model is a probabilistic model that assumes that each cluster follows a Gaussian (or normal) distribution and estimates the weight of the density function for each cluster (Reynolds et al., 2009; Singh et al., 2010). We assume that the sentence vectors of two distinct classes should achieve high accuracy with Gaussian Mixture if they are displayed in a Gaussian distribution in space and are separable from each




Figure 5: Comparison of F1 score per each transformation of best classifiers for sbert and tfidf embedding in no-context setting.




Figure 6: Difference in F1 scores between no-context and contextualized setting as a percentage of contextualized score.

other.

However, there is a potential limitation to using this method. The separability and accuracy score may be underestimated if two clusters are not normally distributed, as illustrated in Figure 7.




Figure 7: Cases that Gaussian Mixture Underestimate the Separability of Two Clusters

Our evaluation results show that out of 78 pairs of derivation classes, 7 pairs achieved an accuracy score above 90%, 24 pairs above 80%, and 47 pairs above 60%, as depicted in Figure 8.

In particular, the class ‘ban’ demonstrates good separability with many other classes, achieving an accuracy score of 0.982 with ‘past’, 0.980 with ‘formal sentence’, 0.942 with ‘minimal change’, and 0.937 with ‘future’, among other pairs.

Additionally, our evaluation results reveal that tenses are generally well-separated, with an accuracy score of 0.90 for ‘past’ and future’ classes, and an accuracy score of 0.924 for ‘simple sentence’ and ‘future’ classes.




Figure 8: Accuracy of Measured Separability with Gaussian Mixture

References

Petra Barančíková and Ondřej Bojar. 2020. Costra 1.1: An inquiry into geometric properties of sentence

spaces. In *Text, Speech, and Dialogue: 23rd International Conference, TSD 2020, Brno, Czech Republic, September 8–11, 2020, Proceedings*, pages 135–143. Springer.

Quoc Le and Tomas Mikolov. 2014. Distributed representations of sentences and documents. In *International conference on machine learning*, pages 1188–1196. PMLR.

Nils Reimers and Iryna Gurevych. 2019. Sentence-bert: Sentence embeddings using siamese bert-networks. *arXiv preprint arXiv:1908.10084*.

Douglas A Reynolds et al. 2009. Gaussian mixture models. *Encyclopedia of biometrics*, 741(659-663).

Ravindra Singh, Bikash C. Pal, and Rabih A. Jabr. 2010. [Statistical representation of distribution system loads using gaussian mixture model](#). *IEEE Transactions on Power Systems*, 25(1):29–37.

A SBERT embeddings

For SBERT implementation we used the `sentence_transformers`² python library.

To produce sentence embeddings with SBERT, we chose pretrained model referred to as ‘*paraphrase-multilingual-MiniLM-L12-v2*’. It offers relatively good scores for its size.

For the supervised SBERT sentence embeddings we list the hyperparameters used in Table 2. Note that due to the small size of training data, the model overfitted quickly. We therefore could afford to train only for 4 epochs.

Hyperparameter	Value
epochs	4
batch size	8
warmup schedule	linear
warmup steps	10000
optimizer	AdamW
weight decay	0.1
learning rate	2e-5

Table 2: Hyperparameters used to train SBERT for supervised embeddings

B Probing method in detail

For probing tasks whose results we described in Section 4.1, we used `scikit-learn`³ python package.

The grid searched parameters and classifiers are presented in Table 3. We used cross-validated grid search, where the performance of a set of parameters is evaluated using cross validation. For all cross validations (both within and outside of grid search) we used `StratifiedGroupKFold`, where seed identifiers were used as groups. This caused the folds to be stratified as well as sharing as few common seed sentences as possible.

We report the best performing set of parameters for each classifier in Tables 4, 5 and 6

²<https://www.sbert.net/>

³<https://scikit-learn.org/stable/>

Classifier	Hyperparameter	Values searched
RandomForestClassifier	n_estimators	50, 100, 200
	max_depth	2, 5, 25, None
	min_samples_split	2, 10, 20
SVC*	kernel	rbf, linear
	gamma	auto, scale
KNeighboursClassifier*	n_neighbours	3, 5, 10
	weights	uniform, distance

Table 3: Classifiers and hyperparameters grid search for probe tasks. All are named equally as in the scikit-learn python package. *Embeddings were first scaled to have zero mean and unit variance.

Context	Embedding	gamma	kernel
diff	doc2vec	auto	rbf
	sbert	auto	rbf
no-context	doc2vec	auto	rbf
	sbert	auto	rbf
	sbert_supervised_0	auto	rbf
	sbert_supervised_2	auto	rbf

Table 4: Best parameters for SVC classifier.

Context	Embedding	max_depth	min_samples_split	n_estimators
diff	bow	25	20	200
	mixup_by_seed_doc2vec	None	5	200
	mixup_by_seed_sbert	25	5	200
	sbert_supervised_0	None	5	200
	sbert_supervised_1	None	5	200
	sbert_supervised_2	25	10	200
	sbert_supervised_3	None	5	200
	sbert_supervised_4	None	10	100
	tfidf	None	20	200
no-context	bow	25	10	200
	mixup_by_seed_doc2vec	None	20	50
	sbert_supervised_1	25	10	50
	sbert_supervised_3	25	10	100
	sbert_supervised_4	25	5	50
	tfidf	25	10	200

Table 5: Best parameters for RandomForestClassifier classifier.

Context	Embedding	n_neighbors	weights
no-context	mixup_by_seed_sbert	5	distance

Table 6: Best parameters for KNeighboursClassifier classifier.